

Distributed Systems

Tutorial #3 Docker Containers

Yaron Hay

Largely based on the slides of
Amir Dachbash

Technion

Israel Institute of Technology
Faculty of Computer Science

Fall 2024/25

November 27, 2024

v1.0 27/11/2024



TECHNION

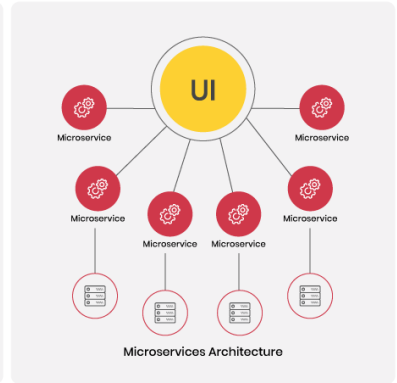



The Henry and Marilyn Taub
Faculty of Computer Science

Motivation

Microservices


- Microservices are self-contained and independent application units
- Each microservice fulfills one specific business function
 - But, still considered as a small application
- What would happen if you create a lots of microservices for your app?
 - Resource scaling would be hard to accomplish
 - Increased deployment complexity
 - But, it worked on my machine!



Key Concept

Docker is a software platform that allows you to **build**, **test**, and **deploy** applications quickly.

- Flexible
 - Lightweight
 - Portable
 - Loosely Coupled
 - Scalable
 - Secure
- Docker packages software into standardized units called containers.
 - Images that are built using docker files, create containers.
 - Images are distributed using registries.
 - Containers include everything the software needs to run e.g. libraries, system tools, code, and runtime.



Key Concept

An image is a **read-only** template with instructions for creating a Docker container.
An image has a designated command to be executed.

You can create a Docker image in one of two ways:

- Interactive Method
 - 1 Run a container from an existing Docker image.
 - 2 Manually change that container environment through a series of live steps.
 - 3 Save the resulting state as a new image.
- Building it from a *Dockerfile*

Docker

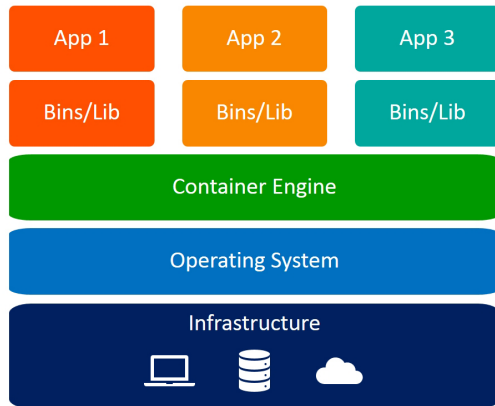
Containers

Key Concept

A container is a **runtime instance of an image**, that *runs until the command completes*.


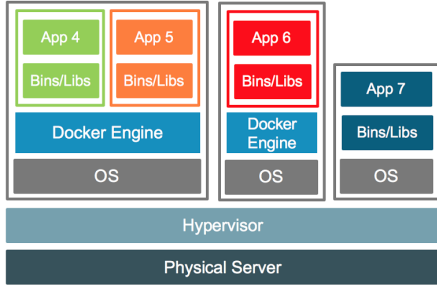
Containers are relatively well isolated from other containers and the host machine, but still lightweight.

- You can create, start, stop, move, or delete a container using the Docker API or CLI.
- You can connect a container to one or more networks, attach storage to it, or even create a new image based on its current state
- The isolation level can be controlled.



Docker

Containers vs Virtual Machines



Docker

Containers vs Virtual Machines

	Virtual Machines	Docker Containers
Underlying Framework	Hypervisor	Docker Engine (directly on host's kernel)
Isolation	Strict	Loose
Typical Size	GBs	MBs
Startup Time	Minutes	Seconds
Use Cases	Production of "heavy systems" e.g, datacenter	Development, testing and distribution of an applications


Table: A comparison between Docker Containers and Virtual Machines



Docker

Layers

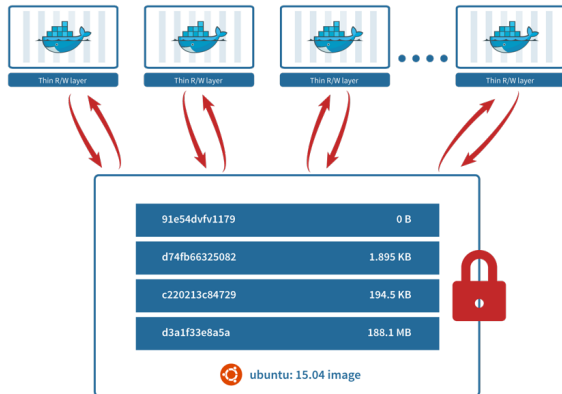
- A Docker image is built up from a series of layers.
- Each layer is **read-only** (except the last one).
- Each layer is only a set of differences from the layer before it.
 - Files are changed in a Copy-on-write manner.
- When you create a new container, a new **writable** layer is added on top of the underlying layers (AKA the “container layer”).
 - All changes made to the running container, such as writing new files, modifying existing files, and deleting files, are written to this thin writable container layer.
 - The files won't persist after the container is deleted



Docker

Layers


- When you the Dockerfile is changed and rebuilt, only the layers that changed are rebuilt.
- Multiple containers can share access to the same underlying image and yet have their own data state.



Docker

Docker Engine

- Docker Engine is installed on the host machine.
- The docker engine follows the client-server architecture.
- The docker engine has 3 components:
 - 1 **The Docker Daemon** called **dockerd** (the server). It can create and manage docker images, containers, networks, etc.
 - 2 **REST API** used to instruct the docker daemon.
 - 3 **Command Line Interface (CLI)** for clients to control the docker daemon.



Key Concept

A stateless, highly scalable server-side application that **stores** and **distributes** Docker images.



- A Docker registry is organized into Docker repositories.
- A repository holds all the versions of a specific image. A version can be accessed by a corresponding *tag*.
- Allows Docker engines can
 - Pull images from the registry (locally)
 - Push new images to the registry
- If no registry is specified, the Docker engine uses with DockerHub, Docker's public registry instance.

Docker

Using the Registry

Interact with a Registry

```
docker pull ubuntu:16.04
docker tag  ubuntu:16.04 myregistry.domain.com/my-ubuntu
docker push  myregistry.domain.com/my-ubuntu
docker pull  myregistry.domain.com/my-ubuntu
```

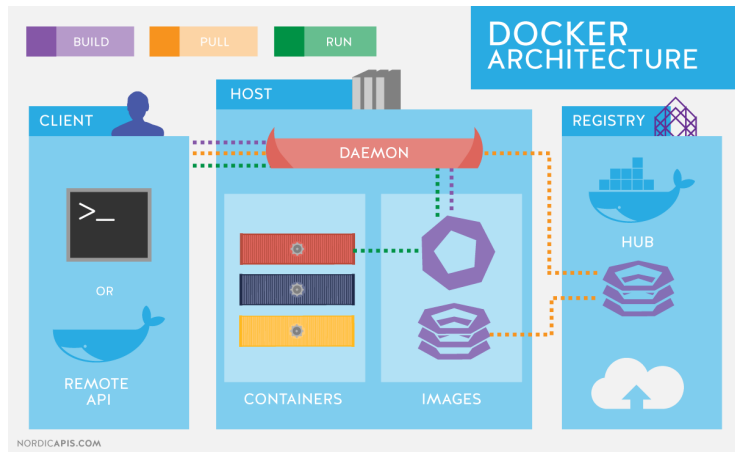
Create your own Registry

```
docker run -d --restart=always --name registry -v "$($pwd)"/certs:/certs \
-e REGISTRY_HTTP_ADDR=0.0.0.0:443 \
-e REGISTRY_HTTP_TLS_CERTIFICATE=/certs/domain.crt -e REGISTRY_HTTP_TLS_KEY=/certs/domain.key \
-p 443:443 registry:2
```

Docker

Architecture


- Developers build images and push them to the registry.
- Clients contact the docker engine via REST API or a CLI.
- Images are pulled from the image registry.
- The docker daemon creates containers from images.



Key Concept

A Dockerfile is a file containing a list of instructions that inform the daemon on **how** to build a docker image.

- Each layer is represented by a **single** instruction in the image's Dockerfile.
- Often, an image is based on another image, with some additional customization.
 - For example, you may build an image which is based on the ubuntu image, but installs the Apache web server and your application



- **FROM** – Defines the base image that is used as the first layer in the build process.
- **RUN** – Executes the command with its arguments in a new layer.
- **ADD** – Copies the files from source on the host to the destination inside the container.
- **CMD** – A command, that gets executed only after the container has been instantiated.
 - **ENTRYPOINT** – Also append the specified command line arguments
- **ENV** – Sets environment variables.
- **WORKDIR** – Sets the working directory for any **RUN**, **CMD**, **ENTRYPOINT**, **COPY** and **ADD** instruction that come after in the Dockerfile.
- **EXPOSE** – Informs Docker that the container listens to the specified network ports during runtime. You can specify whether the container listens for TCP or UDP traffic, (the default is TCP if the protocol is not specified).

Full specification can be found at the [Docker Builder Documentation](#).



Docker Files

Example

```
1 FROM node:10-slim
2
3 RUN apt-get update \
4     && apt-get install -y
5     ↪ --no-install-recommends \
6     curl \
7     && rm -rf /var/lib/apt/lists/*
8
9 ENV TINI_VERSION v0.19.0
10
11 ADD https://github.com/...
12     ↪ .../${TINI_VERSION}/tini /tini
13
14 RUN chmod +x /tini
15
16 WORKDIR /app
```

```
14 RUN npm install -g nodemon
15
16 COPY package*.json ./
17
18 RUN npm ci \
19     && npm cache clean --force \
20     && mv /app/node_modules /node_modules
21
22 COPY . .
23
24 ENV PORT 80
25 EXPOSE 80
26
27 CMD ["/tini", "--", "node", "server.js"]
```



Docker Networking

Network Types


- **Bridge** [The default network driver]
 - Usually used when your applications run in standalone containers that need to communicate with each other
- **Host**
 - Removes the network isolation between the container and the host
 - Allows for directly using the host's networking
- **Overlay**
 - Connects multiple Docker daemons together
 - Enables swarm services to communicate with each other
 - Facilitates communication between two standalone containers on different Docker daemons
- **MACvLAN Networks**
 - Allows you to assign a MAC address to a container, making it appear as a physical device on your network.
 - Traffic is routed to containers by their MAC addresses.
 - Best choice when for legacy applications expecting to be directly connected to the physical network.
- **None** Disable all networking



Storage

Storage Drivers

How does the container access the persistent storage on the host?



Storage

Volumes


- The recommended way to persist data in Docker.
- Stored in a part of the host file-system which is managed by Docker
- Persists even after a container is deleted.
 - `/var/lib/docker/volumes/` on Linux
- Non-Docker processes should not modify this part of the filesystem (Convention).
- They are completely managed by docker through docker CLI or Docker API.
- To start a container with a volume is done by `-v` flag in the run command.
 - `$ docker run v <host dir>:<container dir> ...`
 - `$ docker run v /opt/data:var/lib/mysql mysql`



- Bind Mounts
 - May be stored anywhere on the host system. Non-Docker processes on the Docker host or a Docker container can modify them at any time.
- tmpfs (Linux)
 - Are stored in the host system's memory only
 - Are never written to the host system's filesystem.
 - Why not writing to the 'Container layer'?
 - Performance or security reasons or if the data relates to non-persistent application state.
- Named pipes (Microsoft)



Docker Compose



Source at <https://github.com/dockeramples/example-voting-app>.



TECHNION



The Henry and Marilyn Taub
Faculty of Computer Science

Docker Compose

Example

```
1  services:
2    vote:
3      build: ./vote
4      command: python app.py
5      depends_on:
6        redis:
7          condition: service_healthy
8      volumes:
9        - ./vote:/app
10     ports:
11       - "5000:80"
12     networks:
13       - front-tier
14       - back-tier
15   result:
16     build: ./result
17     # use nodemon rather than node for
18     ↪ local dev
19     command: nodemon server.js
20     depends_on:
21       db:
22         condition: service_healthy
23     volumes:
24       - ./result:/app
25     ports:
26       - "5001:80"
27       - "5858:5858"
28     networks:
29       - front-tier
30       - back-tier
31   worker:
32     build:
33       context: ./worker
34     depends_on:
35       redis:
36         condition: service_healthy
37       db:
38         condition: service_healthy
39     networks:
40       - back-tier
41   redis:
42     image: redis:5.0-alpine3.10
43     volumes:
44       - "./healthchecks:/healthchecks"
45     healthcheck:
46       test: /healthchecks/redis.sh
47       interval: "5s"
48     ports: ["6379"]
49     networks:
50       - back-tier
51   db:
52     image: postgres:9.4
53     environment:
54       POSTGRES_USER: "postgres"
55       POSTGRES_PASSWORD: "postgres"
56     volumes:
57       -
58       ↪ "db-data:/var/lib/postgresql/data"
59       - "./healthchecks:/healthchecks"
60     healthcheck:
61       test: /healthchecks/postgres.sh
62       interval: "5s"
63     networks:
64       - back-tier
65   volumes:
66     db-data:
67     networks:
68       front-tier:
69       back-tier:
```



Why?

Run Containers in a Cluster Mode.

- Redundancy
- High Availability
- Scalability
- Health Monitoring [of containers and hosts]

Further Reading

Docker Swarm <https://docs.docker.com/engine/swarm/key-concepts/>.

Kubernetes <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>.



Container Orchestration

Docker Swarm

- A swarm consists of multiple Docker hosts which run in swarm mode and act as managers and workers.
- A given Docker host can be a manager, a worker, or perform both roles.
- When you create a service, you define its optimal state (number of replicas, network and storage resources available to it, etc..), Docker works to maintain that desired state.
- For instance, if a worker node becomes unavailable, Docker schedules that node's tasks on other nodes.
 - A task is a running container which is part of a swarm service and managed by a swarm manager



Further Reading

- More advanced material <https://docs.docker.com/get-started/overview/>
 - **Namespaces** - provide the isolated workspace called the container.
 - **cgroups** (control groups) - used to limit and isolate the resource usage (CPU, memory, Disk I/O, network, etc.)
 - **UnionFS** (Union file systems)
- Docker and Microservices
 - <https://www.sumologic.com/insight/microservices-architecture-docker-containers/>
 - <https://timber.io/blog/docker-and-the-rise-of-microservices/>
- Running linux containers on windows ?
 - <https://docs.microsoft.com/en-us/virtualization/windowscontainers/deploy-containers/linux-containers>

